

Introduction

MNIST ("Modified National Institute of Standards and Technology") is the de facto "Hello World" dataset of computer vision. Since its release in 1999, this classic dataset of handwritten images has served as the basis for benchmarking classification algorithms. As new machine learning techniques emerge, MNIST remains a reliable resource for researchers and learners alike.

Objective is to correctly identify digits from a dataset of tens of thousands of handwritten images

Approach

For this, we will be using Keras (with TensorFlow as our backend) as the main package to create a simple neural network to predict, as accurately as we can, digits from handwritten images. In particular, we will be calling the Functional Model API of Keras, and creating a 2-layered, 3-layered and 5-layered neural network.

Also, we will be experimenting with various optimizers: the plain vanilla Stochastic Gradient Descent optimizer and the Adam optimizer. However, there are many other parameters, such as training epochs which we will not be experimenting with.

In addition, the choice of hidden layer units are completely arbitrary and may not be optimal. This is yet another parameter which we will not attempt to tinker with. Lastly, we introduce dropout, a form of regularisation, in our neural networks to prevent overfitting.

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
```

```
import seaborn as sns
import matplotlib.pyplot as plt
import tensorflow as tf
print("Printing the version of Tensorflow installed - ", tf.__version__)
```

Printing the version of Tensorflow installed - 1.13.1

Importing the dataset

```
In [2]: mnist_data = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist_data.load_data()


print("Number of training examples :", x_train.shape[0], "and each image is of shape (%d, %d)"%(x_train.shape[1], x_train.shape[2]))
print("Number of training examples :", x_test.shape[0], "and each image is of shape (%d, %d)"%(x_test.shape[1], x_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

Plotting

```
In [3]: print("Printing the label for the first image :", y_train[0])
plt.imshow(x_train[0], cmap=plt.cm.binary)
plt.show()
```

Printing the label for the first image : 5



Data preprocessing and Data cleaning

```
In [4]: x_train.shape
```

```
Out[4]: (60000, 28, 28)
```

Observation - As you can see, the above dataset contains 3D array with each row containing 28x28 matrix. We will have to change this such that each row will contain 784(28*28) columns.

```
In [5]: x_train = x_train.reshape(x_train.shape[0], x_train.shape[1] * x_train.
      shape[2])
x_test = x_test.reshape(x_test.shape[0], x_test.shape[1]*x_test.shape[2]
      ])

print("The shape of Training data now becomes : ", x_train.shape)
print("The shape of Testing data now becomes : ", x_test.shape)
```

```
The shape of Training data now becomes : (60000, 784)
```

```
The shape of Testing data now becomes : (10000, 784)
```

```
In [6]: x_train[0]
```

```
Out[6]: array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  3, 18, 18, 18,
126, 136, 175, 26, 166, 255, 247, 127,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0, 30, 36, 94, 154, 170, 253,
253, 253, 253, 253, 225, 172, 253, 242, 195, 64,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0, 49, 238, 253, 253, 253,
253, 253, 253, 253, 253, 251, 93, 82, 82, 56, 39,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 18, 219, 253,
253, 253, 253, 253, 198, 182, 247, 241,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 80, 156, 107, 253, 253, 205, 11,  0, 43, 154,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0, 14,  1, 154, 253, 90,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0, 139, 253, 190,  2,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 11, 190, 253, 70,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 35,
241, 225, 160, 108,  1,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0, 81, 240, 253, 253, 119, 25,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0, 45, 186, 253, 253, 150, 27,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0, 16, 93, 252, 253, 187,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 249,
```

```

253, 249, 64, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 46, 130,
183, 253, 253, 207, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 39, 148,
229, 253, 253, 253, 250, 182, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 24, 114,
221, 253, 253, 253, 253, 201, 78, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 23, 66,
213, 253, 253, 253, 253, 198, 81, 2, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 18, 171,
219, 253, 253, 253, 253, 195, 80, 9, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 55, 172,
226, 253, 253, 253, 253, 244, 133, 11, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
136, 253, 253, 253, 212, 135, 132, 16, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0], dtype=uint8)

```

Observation - We will have to normalize the data before training the MLP model.

There are two approaches for normalization-

1. Divide the entire dataset by 255.0 (Since in rgb 255 is the maximum value)
2. We can use the inbuilt normalization library of tensorflow - `utils.normalize()`

```

In [0]: x_train = tf.keras.utils.normalize(x_train)
        x_test = tf.keras.utils.normalize(x_test)

```

```

In [8]: # After normalizing
        x_train[0]

```

[illegible]

0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0.02010494, 0.09765254,
0.10380711, 0.10380711, 0.10380711, 0.10380711, 0.10380711,
0.10380711, 0.10380711, 0.10380711, 0.1029865 , 0.03815835,
0.03364499, 0.03364499, 0.02297707, 0.01600189, 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0.00738549, 0.08985675, 0.10380711, 0.10380711,
0.10380711, 0.10380711, 0.10380711, 0.08124035, 0.07467547,
0.10134528, 0.09888346, 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0.03282438, 0.06400755, 0.04390261, 0.10380711, 0.10380711,
0.08411248, 0.00451335, 0. , 0.01764311, 0.06318694,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0.00574427,
0.0004103 , 0.06318694, 0.10380711, 0.03692743, 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0.05703237,
0.10380711, 0.07795791, 0.00082061, 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0.00451335, 0.07795791, 0.10380711,
0.02872134, 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0.01436067, 0.09888346, 0.09231858, 0.06564877,
0.04431292, 0.0004103 , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,

0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0.03323469, 0.09847315, 0.10380711, 0.10380711, 0.04882627,
0.01025762, 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0.01846372,
0.07631669, 0.10380711, 0.10380711, 0.06154572, 0.01107823,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0.00656488, 0.03815835,
0.10339681, 0.10380711, 0.076727 , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0.10216589, 0.10380711,
0.10216589, 0.02625951, 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0.01887402, 0.05333962,
0.07508578, 0.10380711, 0.10380711, 0.08493309, 0.00082061,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0.01600189,
0.06072511, 0.0939598 , 0.10380711, 0.10380711, 0.10380711,
0.1025762 , 0.07467547, 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0. , 0. , 0. , 0. , 0. ,
0.00984732, 0.04677475, 0.09067736, 0.10380711, 0.10380711,
0.10380711, 0.10380711, 0.08247126, 0.03200377, 0. ,
0. , 0. , 0. , 0. , 0. ,

Create PDF in your applications with the Pdfcrowd [HTML to PDF API](#)

```
0.      , 0.      , 0.      , 0.      , 0.      ,  
0.      , 0.      , 0.      , 0.      , 0.      ,
```

One hot encoding

```
In [9]: print("Class label for the first image ", y_train[0])
```

```
y_train = tf.keras.utils.to_categorical(y_train, num_classes=10)  
y_test = tf.keras.utils.to_categorical(y_test, num_classes=10)
```

```
print("Class label after one hot encoding : ", y_train[0])
```

```
Class label for the first image 5
```

```
Class label after one hot encoding : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

MLP Architectures on dataset using Keras

```
In [0]: # Some paramaters to be predefined
```

```
n_epochs = 20  
batchsize = 1024
```

```
output_dim = 10  
input_dim = x_train.shape[1]
```

```
In [0]: final_output = pd.DataFrame(columns=["#Layers", "Model", "Layer-Archite  
cture", "Optimizer", "BN-Present",  
                                           "Dropout-Present",  
                                           "Train-loss", "Test-loss", "Train-  
accuracy", "Test-Accuracy"])
```

```
In [0]: # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4  
# https://stackoverflow.com/a/14434334  
# this function is used to update the plots for each epoch and error  
def plt_dynamic(x, vy, ty, ax, colors=['b']):
```

```

ax.plot(x, vy, 'b', label="Validation Loss")
ax.plot(x, ty, 'r', label="Train Loss")
plt.legend()
plt.grid()
fig.canvas.draw()

```

2 Hidden Layers architecture

2 ReLU hidden Layers (512-128) + ADAM

```

In [13]: relumodel_2 = tf.keras.models.Sequential()
relumodel_2.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input
_shape=(input_dim, )))
relumodel_2.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.soft
max))

relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=b
atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -

```

```

> 128 -> 10",
False,
.history["loss"][n_epochs-1]),
history["val_loss"][n_epochs-1]),
odel.history["acc"][n_epochs-1]),
del.history["val_acc"][n_epochs-1]}}, ignore_index=True)

"Optimizer": "ADAM", "BN-Present":
"Dropout-Present": False,
"Train-loss": '{:.5f}'.format(model
"Test-loss": '{:.5f}'.format(model.
"Train-accuracy": '{:.5f}'.format(m
"Test-Accuracy": '{:.5f}'.format(mo

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Train on 60000 samples, validate on 10000 samples

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/20

60000/60000 [=====] - 4s 63us/sample - loss:

0.9974 - acc: 0.7768 - val_loss: 0.3515 - val_acc: 0.9000

Epoch 2/20

```
60000/60000 [=====] - 4s 59us/sample - loss:
0.3060 - acc: 0.9112 - val_loss: 0.2601 - val_acc: 0.9253
Epoch 3/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.2380 - acc: 0.9311 - val_loss: 0.2112 - val_acc: 0.9393
Epoch 4/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.1939 - acc: 0.9434 - val_loss: 0.1797 - val_acc: 0.9492
Epoch 5/20
60000/60000 [=====] - 4s 59us/sample - loss:
0.1644 - acc: 0.9522 - val_loss: 0.1556 - val_acc: 0.9533
Epoch 6/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.1415 - acc: 0.9592 - val_loss: 0.1402 - val_acc: 0.9573
Epoch 7/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.1224 - acc: 0.9644 - val_loss: 0.1242 - val_acc: 0.9629
Epoch 8/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.1080 - acc: 0.9688 - val_loss: 0.1127 - val_acc: 0.9663
Epoch 9/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0950 - acc: 0.9728 - val_loss: 0.1033 - val_acc: 0.9699
Epoch 10/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0848 - acc: 0.9748 - val_loss: 0.0965 - val_acc: 0.9710
Epoch 11/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0754 - acc: 0.9787 - val_loss: 0.0897 - val_acc: 0.9727
Epoch 12/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0678 - acc: 0.9804 - val_loss: 0.0883 - val_acc: 0.9738
Epoch 13/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0607 - acc: 0.9823 - val_loss: 0.0831 - val_acc: 0.9752
Epoch 14/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0546 - acc: 0.9844 - val_loss: 0.0781 - val_acc: 0.9763
Epoch 15/20
```

```

60000/60000 [=====] - 4s 61us/sample - loss:
0.0497 - acc: 0.9861 - val_loss: 0.0786 - val_acc: 0.9758
Epoch 16/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0449 - acc: 0.9876 - val_loss: 0.0745 - val_acc: 0.9773
Epoch 17/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0407 - acc: 0.9887 - val_loss: 0.0711 - val_acc: 0.9786
Epoch 18/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0375 - acc: 0.9899 - val_loss: 0.0722 - val_acc: 0.9779
Epoch 19/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0341 - acc: 0.9911 - val_loss: 0.0690 - val_acc: 0.9787
Epoch 20/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0306 - acc: 0.9920 - val_loss: 0.0699 - val_acc: 0.9784
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dense_1 (Dense)	(None, 128)	65664
dense_2 (Dense)	(None, 10)	1290

```

=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

```

None

```


*****

```

```

10000/10000 [=====] - 1s 74us/sample - loss:
0.0699 - acc: 0.9784
Test score: 0.06986413442818448
Test accuracy: 0.9784

```



```
In [14]: w_after = relumodel_2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ' )  
plt.show()
```



2 ReLU hidden Layers (256-256) + ADAM

```
In [15]: relumodel_2 = tf.keras.models.Sequential()  
relumodel_2.add(tf.keras.layers.Dense(256, activation=tf.nn.relu, input_shape=(input_dim, )))  
relumodel_2.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))  
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))  
  
relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy',  
metrics=['accuracy'])  
  
model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))
```



```

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 256 -
> 256 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.
history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 45us/sample - loss:

```
1.1184 - acc: 0.7268 - val_loss: 0.3768 - val_acc: 0.8929
Epoch 2/20
60000/60000 [=====] - 3s 42us/sample - loss:
0.3240 - acc: 0.9049 - val_loss: 0.2743 - val_acc: 0.9183
Epoch 3/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.2565 - acc: 0.9250 - val_loss: 0.2287 - val_acc: 0.9324
Epoch 4/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.2157 - acc: 0.9383 - val_loss: 0.1954 - val_acc: 0.9418
Epoch 5/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.1863 - acc: 0.9456 - val_loss: 0.1679 - val_acc: 0.9506
Epoch 6/20
60000/60000 [=====] - 2s 40us/sample - loss:
0.1610 - acc: 0.9529 - val_loss: 0.1536 - val_acc: 0.9540
Epoch 7/20
60000/60000 [=====] - 2s 39us/sample - loss:
0.1419 - acc: 0.9582 - val_loss: 0.1375 - val_acc: 0.9579
Epoch 8/20
60000/60000 [=====] - 2s 39us/sample - loss:
0.1245 - acc: 0.9635 - val_loss: 0.1254 - val_acc: 0.9609
Epoch 9/20
60000/60000 [=====] - 2s 40us/sample - loss:
0.1110 - acc: 0.9676 - val_loss: 0.1161 - val_acc: 0.9634
Epoch 10/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.1000 - acc: 0.9708 - val_loss: 0.1085 - val_acc: 0.9659
Epoch 11/20
60000/60000 [=====] - 2s 40us/sample - loss:
0.0897 - acc: 0.9741 - val_loss: 0.1051 - val_acc: 0.9673
Epoch 12/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.0816 - acc: 0.9756 - val_loss: 0.0963 - val_acc: 0.9697
Epoch 13/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.0737 - acc: 0.9786 - val_loss: 0.0928 - val_acc: 0.9710

Epoch 14/20
60000/60000 [=====] - 2s 41us/sample - loss:
```

```

0.0684 - acc: 0.9800 - val_loss: 0.0870 - val_acc: 0.9718
Epoch 15/20
60000/60000 [=====] - 3s 42us/sample - loss:
0.0621 - acc: 0.9819 - val_loss: 0.0862 - val_acc: 0.9734
Epoch 16/20
60000/60000 [=====] - 3s 42us/sample - loss:
0.0561 - acc: 0.9839 - val_loss: 0.0824 - val_acc: 0.9748
Epoch 17/20
60000/60000 [=====] - 2s 42us/sample - loss:
0.0512 - acc: 0.9855 - val_loss: 0.0806 - val_acc: 0.9742
Epoch 18/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.0470 - acc: 0.9866 - val_loss: 0.0800 - val_acc: 0.9743
Epoch 19/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.0447 - acc: 0.9874 - val_loss: 0.0762 - val_acc: 0.9759
Epoch 20/20
60000/60000 [=====] - 2s 41us/sample - loss:
0.0402 - acc: 0.9887 - val_loss: 0.0780 - val_acc: 0.9763

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 256)	200960
dense_4 (Dense)	(None, 256)	65792
dense_5 (Dense)	(None, 10)	2570

```

Total params: 269,322
Trainable params: 269,322
Non-trainable params: 0

```

None

```

10000/10000 [=====] - 1s 51us/sample - loss:

```

```

0.0780 - acc: 0.9763


```

```

Test score: 0.077971170845069

```

Test accuracy: 0.9763



```
In [16]: w_after = relumodel_2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 ReLU hidden Layers (384-128) + ADAM

```
In [17]: relumodel_2 = tf.keras.models.Sequential()
relumodel_2.add(tf.keras.layers.Dense(384, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_2.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))
```

```

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 384 -
> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20

```
60000/60000 [=====] - 3s 49us/sample - loss:
1.0646 - acc: 0.7678 - val_loss: 0.3695 - val_acc: 0.8988
Epoch 2/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.3195 - acc: 0.9073 - val_loss: 0.2722 - val_acc: 0.9203
Epoch 3/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.2506 - acc: 0.9271 - val_loss: 0.2243 - val_acc: 0.9339
Epoch 4/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.2089 - acc: 0.9398 - val_loss: 0.1871 - val_acc: 0.9460
Epoch 5/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.1766 - acc: 0.9483 - val_loss: 0.1664 - val_acc: 0.9527
Epoch 6/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.1520 - acc: 0.9552 - val_loss: 0.1488 - val_acc: 0.9554
Epoch 7/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.1332 - acc: 0.9612 - val_loss: 0.1313 - val_acc: 0.9620
Epoch 8/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.1172 - acc: 0.9660 - val_loss: 0.1229 - val_acc: 0.9636
Epoch 9/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.1051 - acc: 0.9694 - val_loss: 0.1138 - val_acc: 0.9650
Epoch 10/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.0933 - acc: 0.9727 - val_loss: 0.1066 - val_acc: 0.9677
Epoch 11/20
60000/60000 [=====] - 3s 47us/sample - loss:
0.0838 - acc: 0.9754 - val_loss: 0.0974 - val_acc: 0.9696
Epoch 12/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0748 - acc: 0.9783 - val_loss: 0.0913 - val_acc: 0.9713
Epoch 13/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0678 - acc: 0.9806 - val_loss: 0.0870 - val_acc: 0.9732

Epoch 14/20
```

```

60000/60000 [=====] - 3s 46us/sample - loss:
0.0618 - acc: 0.9825 - val_loss: 0.0852 - val_acc: 0.9735
Epoch 15/20
60000/60000 [=====] - 3s 45us/sample - loss:
0.0551 - acc: 0.9846 - val_loss: 0.0815 - val_acc: 0.9745
Epoch 16/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0508 - acc: 0.9857 - val_loss: 0.0771 - val_acc: 0.9759
Epoch 17/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0460 - acc: 0.9874 - val_loss: 0.0747 - val_acc: 0.9760
Epoch 18/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0415 - acc: 0.9890 - val_loss: 0.0741 - val_acc: 0.9786
Epoch 19/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0378 - acc: 0.9899 - val_loss: 0.0740 - val_acc: 0.9769
Epoch 20/20
60000/60000 [=====] - 3s 46us/sample - loss:
0.0355 - acc: 0.9904 - val_loss: 0.0729 - val_acc: 0.9767

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
=====		
dense_6 (Dense)	(None, 384)	301440
dense_7 (Dense)	(None, 128)	49280
dense_8 (Dense)	(None, 10)	1290
=====		

Total params: 352,010
Trainable params: 352,010
Non-trainable params: 0

None


```

10000/10000 [=====] - 1s 57us/sample - loss:
0.0729 - acc: 0.9767

```


Test score: 0.07292742731682957

Test accuracy: 0.9767



```
In [18]: w_after = relumodel_2.get_weights()

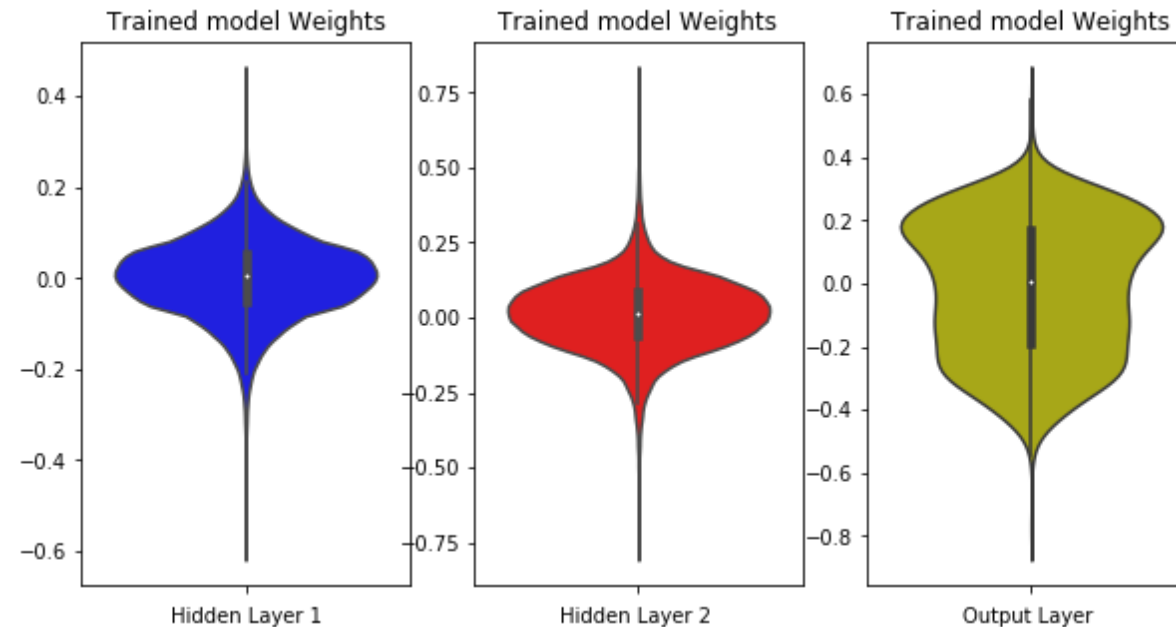
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 ReLU hidden Layers (512-128) + BatchNormalization + Dropout + ADAM

```
In [19]: relumodel_2 = tf.keras.models.Sequential()
relumodel_2.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))
```

```

relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=b
atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -
> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

```

```
vy = model.history['val_loss']  
ty = model.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/layers/core.py:143: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 87us/sample - loss: 0.6757 - acc: 0.7988 - val_loss: 1.8734 - val_acc: 0.3371

Epoch 2/20

60000/60000 [=====] - 5s 75us/sample - loss: 0.2986 - acc: 0.9108 - val_loss: 1.7556 - val_acc: 0.2403

Epoch 3/20

60000/60000 [=====] - 5s 76us/sample - loss: 0.2363 - acc: 0.9294 - val_loss: 1.5945 - val_acc: 0.2704

Epoch 4/20

60000/60000 [=====] - 5s 77us/sample - loss: 0.1938 - acc: 0.9416 - val_loss: 1.3605 - val_acc: 0.4257

Epoch 5/20

60000/60000 [=====] - 5s 77us/sample - loss: 0.1655 - acc: 0.9500 - val_loss: 1.0498 - val_acc: 0.7379

Epoch 6/20

60000/60000 [=====] - 5s 77us/sample - loss: 0.1466 - acc: 0.9566 - val_loss: 0.6820 - val_acc: 0.9064

Epoch 7/20

60000/60000 [=====] - 5s 78us/sample - loss: 0.1305 - acc: 0.9608 - val_loss: 0.3947 - val_acc: 0.9472

Epoch 8/20

60000/60000 [=====] - 5s 78us/sample - loss: 0.1169 - acc: 0.9643 - val_loss: 0.2245 - val_acc: 0.9594

Epoch 9/20

60000/60000 [=====] - 5s 79us/sample - loss: 0.1070 - acc: 0.9673 - val_loss: 0.1303 - val_acc: 0.9742

```

Epoch 10/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0971 - acc: 0.9697 - val_loss: 0.0888 - val_acc: 0.9762
Epoch 11/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0890 - acc: 0.9731 - val_loss: 0.0727 - val_acc: 0.9778
Epoch 12/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0843 - acc: 0.9743 - val_loss: 0.0689 - val_acc: 0.9776
Epoch 13/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0774 - acc: 0.9757 - val_loss: 0.0665 - val_acc: 0.9796
Epoch 14/20
60000/60000 [=====] - 5s 79us/sample - loss:
0.0741 - acc: 0.9767 - val_loss: 0.0651 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0708 - acc: 0.9776 - val_loss: 0.0658 - val_acc: 0.9805
Epoch 16/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0665 - acc: 0.9790 - val_loss: 0.0623 - val_acc: 0.9817
Epoch 17/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0616 - acc: 0.9800 - val_loss: 0.0615 - val_acc: 0.9811
Epoch 18/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0579 - acc: 0.9816 - val_loss: 0.0604 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 5s 75us/sample - loss:
0.0549 - acc: 0.9827 - val_loss: 0.0618 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 5s 75us/sample - loss:
0.0516 - acc: 0.9832 - val_loss: 0.0597 - val_acc: 0.9823
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 512)	401920

batch_normalization_v1 (Batch Normalization)	(None, 512)	2048
dropout (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 128)	65664
batch_normalization_v1_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		


None

10000/10000 [=====] - 1s 72us/sample - loss:

0.0597 - acc: 0.9823

Test score: 0.05972711388359894

Test accuracy: 0.9823




```
In [20]: w_after = relumodel_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 ReLU hidden Layers (256-256) + BatchNormalization + Dropout + ADAM

```
In [21]: relumodel_2 = tf.keras.models.Sequential()
relumodel_2.add(tf.keras.layers.Dense(256, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=b
```



```

atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 256 -
> 256 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 4s 68us/sample - loss:
0.7771 - acc: 0.7660 - val_loss: 1.9486 - val_acc: 0.1833
Epoch 2/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.3251 - acc: 0.9003 - val_loss: 1.8251 - val_acc: 0.1206
Epoch 3/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.2563 - acc: 0.9229 - val_loss: 1.6905 - val_acc: 0.1385
Epoch 4/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.2158 - acc: 0.9355 - val_loss: 1.4314 - val_acc: 0.2957
Epoch 5/20
60000/60000 [=====] - 4s 64us/sample - loss:
0.1894 - acc: 0.9422 - val_loss: 1.0783 - val_acc: 0.6259
Epoch 6/20
60000/60000 [=====] - 4s 65us/sample - loss:
0.1683 - acc: 0.9496 - val_loss: 0.7199 - val_acc: 0.8310
Epoch 7/20
60000/60000 [=====] - 4s 65us/sample - loss:
0.1531 - acc: 0.9531 - val_loss: 0.4377 - val_acc: 0.9161
Epoch 8/20
60000/60000 [=====] - 4s 66us/sample - loss:
0.1384 - acc: 0.9571 - val_loss: 0.2402 - val_acc: 0.9458
Epoch 9/20
60000/60000 [=====] - 4s 65us/sample - loss:
0.1287 - acc: 0.9596 - val_loss: 0.1465 - val_acc: 0.9665
Epoch 10/20
60000/60000 [=====] - 4s 64us/sample - loss:
0.1168 - acc: 0.9634 - val_loss: 0.1015 - val_acc: 0.9724
Epoch 11/20
60000/60000 [=====] - 4s 59us/sample - loss:
0.1117 - acc: 0.9648 - val_loss: 0.0820 - val_acc: 0.9751
Epoch 12/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.1063 - acc: 0.9669 - val_loss: 0.0749 - val_acc: 0.9764
Epoch 13/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0998 - acc: 0.9680 - val_loss: 0.0724 - val_acc: 0.9772
```

```

Epoch 14/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0946 - acc: 0.9709 - val_loss: 0.0693 - val_acc: 0.9784
Epoch 15/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0894 - acc: 0.9724 - val_loss: 0.0677 - val_acc: 0.9790
Epoch 16/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0882 - acc: 0.9728 - val_loss: 0.0688 - val_acc: 0.9784
Epoch 17/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0823 - acc: 0.9732 - val_loss: 0.0683 - val_acc: 0.9783
Epoch 18/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0785 - acc: 0.9745 - val_loss: 0.0672 - val_acc: 0.9788
Epoch 19/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0748 - acc: 0.9762 - val_loss: 0.0683 - val_acc: 0.9796
Epoch 20/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0751 - acc: 0.9757 - val_loss: 0.0664 - val_acc: 0.9802
*****

```

Printing the Model Summary


Layer (type)	Output Shape	Param #
=====		
dense_12 (Dense)	(None, 256)	200960
batch_normalization_v1_2 (Batch Normalization)	(None, 256)	1024
dropout_2 (Dropout)	(None, 256)	0
dense_13 (Dense)	(None, 256)	65792
batch_normalization_v1_3 (Batch Normalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
=====		
dense_14 (Dense)	(None, 10)	2570

```
=====
Total params: 271,370
Trainable params: 270,346
Non-trainable params: 1,024
```

None

```
*****
```

```
10000/10000 [=====] - 1s 69us/sample - loss:
0.0664 - acc: 0.9802
Test score: 0.06643414922667434
Test accuracy: 0.9802
```



```
In [22]: w_after = relumodel_2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2 ReLU hidden Layers (384-128) + BatchNormalization + Dropout + ADAM

```

In [23]: relumodel_2 = tf.keras.models.Sequential()
relumodel_2.add(tf.keras.layers.Dense(384, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_2.add(tf.keras.layers.BatchNormalization())
relumodel_2.add(tf.keras.layers.Dropout(0.5))
relumodel_2.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_2.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_2.summary())
print("*****")

score = relumodel_2.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 2,
                                     "Model": "2-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 384 -> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m

```

```

odel.history["acc"][n_epochs-1]),
                                "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 75us/sample - loss: 0.7401 - acc: 0.7775 - val_loss: 1.8800 - val_acc: 0.5990

Epoch 2/20

60000/60000 [=====] - 4s 67us/sample - loss: 0.3188 - acc: 0.9033 - val_loss: 1.6718 - val_acc: 0.6008

Epoch 3/20

60000/60000 [=====] - 4s 74us/sample - loss: 0.2496 - acc: 0.9271 - val_loss: 1.4658 - val_acc: 0.7006

Epoch 4/20

60000/60000 [=====] - 4s 69us/sample - loss: 0.2086 - acc: 0.9377 - val_loss: 1.2000 - val_acc: 0.8156

Epoch 5/20

60000/60000 [=====] - 4s 65us/sample - loss: 0.1793 - acc: 0.9449 - val_loss: 0.8920 - val_acc: 0.9021

Epoch 6/20

60000/60000 [=====] - 4s 65us/sample - loss: 0.1600 - acc: 0.9522 - val_loss: 0.5971 - val_acc: 0.9418

Epoch 7/20

60000/60000 [=====] - 4s 66us/sample - loss: 0.1432 - acc: 0.9571 - val_loss: 0.3685 - val_acc: 0.9577

Epoch 8/20

60000/60000 [=====] - 4s 63us/sample - loss: 0.1290 - acc: 0.9616 - val_loss: 0.2124 - val_acc: 0.9699

```
Epoch 9/20
60000/60000 [=====] - 4s 64us/sample - loss:
0.1195 - acc: 0.9630 - val_loss: 0.1309 - val_acc: 0.9731
Epoch 10/20
60000/60000 [=====] - 4s 65us/sample - loss:
0.1077 - acc: 0.9676 - val_loss: 0.0941 - val_acc: 0.9756
Epoch 11/20
60000/60000 [=====] - 4s 65us/sample - loss:
0.0997 - acc: 0.9694 - val_loss: 0.0809 - val_acc: 0.9760
Epoch 12/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0945 - acc: 0.9711 - val_loss: 0.0747 - val_acc: 0.9770
Epoch 13/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0897 - acc: 0.9724 - val_loss: 0.0721 - val_acc: 0.9768
Epoch 14/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0828 - acc: 0.9742 - val_loss: 0.0715 - val_acc: 0.9775
Epoch 15/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0782 - acc: 0.9751 - val_loss: 0.0701 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0775 - acc: 0.9755 - val_loss: 0.0692 - val_acc: 0.9788
Epoch 17/20
60000/60000 [=====] - 4s 64us/sample - loss:
0.0735 - acc: 0.9767 - val_loss: 0.0652 - val_acc: 0.9799
Epoch 18/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0693 - acc: 0.9782 - val_loss: 0.0698 - val_acc: 0.9793
Epoch 19/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0661 - acc: 0.9789 - val_loss: 0.0679 - val_acc: 0.9792
Epoch 20/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0631 - acc: 0.9801 - val_loss: 0.0662 - val_acc: 0.9803
*****
```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 384)	301440
batch_normalization_v1_4 (Ba	(None, 384)	1536
dropout_4 (Dropout)	(None, 384)	0
dense_16 (Dense)	(None, 128)	49280
batch_normalization_v1_5 (Ba	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_17 (Dense)	(None, 10)	1290

=====
 Total params: 354,058
 Trainable params: 353,034
 Non-trainable params: 1,024
 =====


None

10000/10000 [=====] - 1s 69us/sample - loss:

0.0662 - acc: 0.9803

Test score: 0.06616050415177015

Test accuracy: 0.9803



```
In [24]: w_after = relumodel_2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')
plt.show()
```



3 Hidden Layers architecture

```
In [0]: final_output = final_output.append({"#Layers": "--",
                                             "Model": "--",
                                             "Layer-Architecture": "--",
                                             "Optimizer": "--", "BN-Present": "-
                                             -",
                                             "Dropout-Present": "--",
                                             "Train-loss": "--",
                                             "Test-loss": "--",
                                             "Train-accuracy": "--",
                                             "Test-Accuracy": "--"}, ignore_inde
                                             x=True)
```

3 ReLU hidden Layers (512-256-128) + ADAM

```
In [26]: relumodel_3 = tf.keras.models.Sequential()
relumodel_3.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_3.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 3,
                                     "Model": "3-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 256 -> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1])})
```

```

                                "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1]), ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 79us/sample - loss: 0.8634 - acc: 0.7693 - val_loss: 0.2947 - val_acc: 0.9133

Epoch 2/20

60000/60000 [=====] - 4s 73us/sample - loss: 0.2597 - acc: 0.9237 - val_loss: 0.2218 - val_acc: 0.9331

Epoch 3/20

60000/60000 [=====] - 4s 72us/sample - loss: 0.1966 - acc: 0.9419 - val_loss: 0.1723 - val_acc: 0.9481

Epoch 4/20

60000/60000 [=====] - 4s 73us/sample - loss: 0.1584 - acc: 0.9532 - val_loss: 0.1499 - val_acc: 0.9547

Epoch 5/20

60000/60000 [=====] - 4s 74us/sample - loss: 0.1279 - acc: 0.9622 - val_loss: 0.1265 - val_acc: 0.9615

Epoch 6/20

60000/60000 [=====] - 4s 73us/sample - loss: 0.1064 - acc: 0.9686 - val_loss: 0.1099 - val_acc: 0.9656

Epoch 7/20

60000/60000 [=====] - 4s 74us/sample - loss: 0.0922 - acc: 0.9725 - val_loss: 0.0996 - val_acc: 0.9692

Epoch 8/20

60000/60000 [=====] - 4s 73us/sample - loss: 0.0765 - acc: 0.9771 - val_loss: 0.0919 - val_acc: 0.9717

Epoch 9/20

```

Epoch 9/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0672 - acc: 0.9800 - val_loss: 0.0857 - val_acc: 0.9726
Epoch 10/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0580 - acc: 0.9824 - val_loss: 0.0825 - val_acc: 0.9735
Epoch 11/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0489 - acc: 0.9856 - val_loss: 0.0755 - val_acc: 0.9758
Epoch 12/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0410 - acc: 0.9885 - val_loss: 0.0772 - val_acc: 0.9752
Epoch 13/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0365 - acc: 0.9892 - val_loss: 0.0720 - val_acc: 0.9774
Epoch 14/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0309 - acc: 0.9913 - val_loss: 0.0719 - val_acc: 0.9787
Epoch 15/20
60000/60000 [=====] - 4s 72us/sample - loss:
0.0260 - acc: 0.9930 - val_loss: 0.0733 - val_acc: 0.9773
Epoch 16/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0221 - acc: 0.9942 - val_loss: 0.0674 - val_acc: 0.9797
Epoch 17/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0187 - acc: 0.9954 - val_loss: 0.0687 - val_acc: 0.9788
Epoch 18/20
60000/60000 [=====] - 4s 72us/sample - loss:
0.0160 - acc: 0.9960 - val_loss: 0.0704 - val_acc: 0.9783
Epoch 19/20
60000/60000 [=====] - 4s 72us/sample - loss:
0.0137 - acc: 0.9967 - val_loss: 0.0708 - val_acc: 0.9786
Epoch 20/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0124 - acc: 0.9971 - val_loss: 0.0729 - val_acc: 0.9782
*****

```

Printing the Model Summary


Layer (type)	Output Shape	Param #
--------------	--------------	---------

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 512)	401920
dense_19 (Dense)	(None, 256)	131328
dense_20 (Dense)	(None, 128)	32896
dense_21 (Dense)	(None, 10)	1290

Total params: 567,434
 Trainable params: 567,434
 Non-trainable params: 0

None

10000/10000 [=====] - 1s 89us/sample - loss:
 0.0729 - acc: 0.9782
 Test score: 0.07288095771700609
 Test accuracy: 0.9782




```
In [27]: w_after = relu_model_3.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

3 ReLU hidden Layers (512-128-64) + ADAM

```
In [28]: relumodel_3 = tf.keras.models.Sequential()
relumodel_3.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
```

```

print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 3,
                                     "Model": "3-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -
> 128 -> 64 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 68us/sample - loss:
0.9820 - acc: 0.7609 - val_loss: 0.3465 - val_acc: 0.8991

```
Epoch 2/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.2937 - acc: 0.9139 - val_loss: 0.2482 - val_acc: 0.9270
Epoch 3/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.2216 - acc: 0.9358 - val_loss: 0.1922 - val_acc: 0.9430
Epoch 4/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.1764 - acc: 0.9481 - val_loss: 0.1647 - val_acc: 0.9502
Epoch 5/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.1438 - acc: 0.9576 - val_loss: 0.1383 - val_acc: 0.9591
Epoch 6/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.1215 - acc: 0.9636 - val_loss: 0.1322 - val_acc: 0.9607
Epoch 7/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.1036 - acc: 0.9696 - val_loss: 0.1066 - val_acc: 0.9683
Epoch 8/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0900 - acc: 0.9732 - val_loss: 0.0995 - val_acc: 0.9696
Epoch 9/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0790 - acc: 0.9765 - val_loss: 0.0988 - val_acc: 0.9698
Epoch 10/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0699 - acc: 0.9793 - val_loss: 0.0903 - val_acc: 0.9729
Epoch 11/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0611 - acc: 0.9824 - val_loss: 0.0875 - val_acc: 0.9728
Epoch 12/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0551 - acc: 0.9837 - val_loss: 0.0805 - val_acc: 0.9757
Epoch 13/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0494 - acc: 0.9860 - val_loss: 0.0763 - val_acc: 0.9776
Epoch 14/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0437 - acc: 0.9875 - val_loss: 0.0767 - val_acc: 0.9778
```

```

Epoch 15/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0381 - acc: 0.9895 - val_loss: 0.0899 - val_acc: 0.9733
Epoch 16/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0378 - acc: 0.9888 - val_loss: 0.0797 - val_acc: 0.9772
Epoch 17/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0300 - acc: 0.9918 - val_loss: 0.0728 - val_acc: 0.9790
Epoch 18/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0270 - acc: 0.9929 - val_loss: 0.0689 - val_acc: 0.9802
Epoch 19/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0235 - acc: 0.9939 - val_loss: 0.0766 - val_acc: 0.9780
Epoch 20/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0215 - acc: 0.9948 - val_loss: 0.0732 - val_acc: 0.9800
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 512)	401920
dense_23 (Dense)	(None, 128)	65664
dense_24 (Dense)	(None, 64)	8256
dense_25 (Dense)	(None, 10)	650

```

=====
Total params: 476,490
Trainable params: 476,490
Non-trainable params: 0

```

None


```

10000/10000 [=====] - 1s 69us/sample - loss:
0.0732 - acc: 0.9800

```

Test score: 0.07322663756180554

Test accuracy: 0.98




```
In [29]: w_after = relumodel_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 ReLU hidden Layers (384-256-128) + ADAM

```
In [30]: relumodel_3 = tf.keras.models.Sequential()
relumodel_3.add(tf.keras.layers.Dense(384, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_3.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({
    "#Layers": 3,
    "Model": "3-ReLU + Softmax",
    "Layer-Architecture": "784 -> 384 -> 256 -> 128 -> 10",
    "Optimizer": "ADAM", "BN-Present": False,
    "Dropout-Present": False,
    "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
    "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
    "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1]),
    "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, n_epochs+1))

vy = model.history['val_loss']

```

```
ty = model.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 66us/sample - loss: 0.9044 - acc: 0.7855 - val_loss: 0.3027 - val_acc: 0.9101

Epoch 2/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.2669 - acc: 0.9223 - val_loss: 0.2134 - val_acc: 0.9357

Epoch 3/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.1952 - acc: 0.9424 - val_loss: 0.1727 - val_acc: 0.9489

Epoch 4/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.1569 - acc: 0.9532 - val_loss: 0.1476 - val_acc: 0.9549

Epoch 5/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.1287 - acc: 0.9616 - val_loss: 0.1300 - val_acc: 0.9606

Epoch 6/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.1081 - acc: 0.9675 - val_loss: 0.1110 - val_acc: 0.9661

Epoch 7/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.0918 - acc: 0.9719 - val_loss: 0.1061 - val_acc: 0.9662

Epoch 8/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.0800 - acc: 0.9759 - val_loss: 0.0916 - val_acc: 0.9708

Epoch 9/20

60000/60000 [=====] - 4s 59us/sample - loss: 0.0666 - acc: 0.9803 - val_loss: 0.0872 - val_acc: 0.9723

Epoch 10/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.0592 - acc: 0.9823 - val_loss: 0.0897 - val_acc: 0.9716

Epoch 11/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.0528 - acc: 0.9844 - val_loss: 0.0900 - val_acc: 0.9711

Epoch 12/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.0442 - acc: 0.9875 - val_loss: 0.0752 - val_acc: 0.9757

Epoch 13/20

60000/60000 [=====] - 4s 60us/sample - loss:
0.0387 - acc: 0.9890 - val_loss: 0.0743 - val_acc: 0.9764

Epoch 14/20

60000/60000 [=====] - 4s 60us/sample - loss:
0.0334 - acc: 0.9907 - val_loss: 0.0743 - val_acc: 0.9772

Epoch 15/20

60000/60000 [=====] - 4s 59us/sample - loss:
0.0293 - acc: 0.9916 - val_loss: 0.0702 - val_acc: 0.9795

Epoch 16/20

60000/60000 [=====] - 4s 59us/sample - loss:
0.0253 - acc: 0.9935 - val_loss: 0.0746 - val_acc: 0.9771

Epoch 17/20

60000/60000 [=====] - 4s 62us/sample - loss:
0.0228 - acc: 0.9940 - val_loss: 0.0681 - val_acc: 0.9792

Epoch 18/20

60000/60000 [=====] - 4s 69us/sample - loss:
0.0198 - acc: 0.9948 - val_loss: 0.0751 - val_acc: 0.9770

Epoch 19/20

60000/60000 [=====] - 4s 69us/sample - loss:
0.0161 - acc: 0.9961 - val_loss: 0.0697 - val_acc: 0.9787

Epoch 20/20

60000/60000 [=====] - 4s 58us/sample - loss:
0.0130 - acc: 0.9971 - val_loss: 0.0717 - val_acc: 0.9786

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 384)	301440
dense_27 (Dense)	(None, 256)	98560
dense_28 (Dense)	(None, 128)	32896
dense_29 (Dense)	(None, 10)	1290

Total params: 434,186

Trainable params: 434,186

Non trainable params: 0

non-trainable params: 0


None

10000/10000 [=====] - 1s 69us/sample - loss:

0.0717 - acc: 0.9786

Test score: 0.07172299538475926

Test accuracy: 0.9786




```
In [31]: w_after = relumodel_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 ReLU hidden Layers (512-256-128) + BatchNormalization + Dropout + ADAM

```
In [32]: relumodel_3 = tf.keras.models.Sequential()
relumodel_3.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
```

```

relumodel_3.add(tf.keras.layers.BatchNormalization())
relumodel_3.add(tf.keras.layers.Dropout(0.5))
relumodel_3.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.BatchNormalization())
relumodel_3.add(tf.keras.layers.Dropout(0.5))
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 3,
                                     "Model": "3-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 256 -> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(model.evaluate(x_test, y_test)[1])})

```

```
del.history["val_acc"][n_epochs-1]}), ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 109us/sample - loss: 0.6203 - acc: 0.8054 - val_loss: 1.9858 - val_acc: 0.0978

Epoch 2/20

60000/60000 [=====] - 6s 96us/sample - loss: 0.2467 - acc: 0.9247 - val_loss: 2.3146 - val_acc: 0.0974

Epoch 3/20

60000/60000 [=====] - 6s 95us/sample - loss: 0.1933 - acc: 0.9406 - val_loss: 2.4727 - val_acc: 0.0979

Epoch 4/20

60000/60000 [=====] - 6s 94us/sample - loss: 0.1586 - acc: 0.9510 - val_loss: 2.2506 - val_acc: 0.1620

Epoch 5/20

60000/60000 [=====] - 6s 95us/sample - loss: 0.1383 - acc: 0.9575 - val_loss: 1.9394 - val_acc: 0.1974

Epoch 6/20

60000/60000 [=====] - 6s 96us/sample - loss: 0.1181 - acc: 0.9625 - val_loss: 1.2399 - val_acc: 0.4169

Epoch 7/20

60000/60000 [=====] - 6s 95us/sample - loss: 0.1099 - acc: 0.9666 - val_loss: 0.6685 - val_acc: 0.7272

Epoch 8/20

60000/60000 [=====] - 6s 96us/sample - loss: 0.0987 - acc: 0.9695 - val_loss: 0.2438 - val_acc: 0.9250

Epoch 9/20

60000/60000 [=====] - 6s 95us/sample - loss:

```

0.0898 - acc: 0.9711 - val_loss: 0.1340 - val_acc: 0.9588
Epoch 10/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0809 - acc: 0.9738 - val_loss: 0.0859 - val_acc: 0.9711
Epoch 11/20
60000/60000 [=====] - 6s 96us/sample - loss:
0.0764 - acc: 0.9759 - val_loss: 0.0698 - val_acc: 0.9781
Epoch 12/20
60000/60000 [=====] - 6s 96us/sample - loss:
0.0725 - acc: 0.9766 - val_loss: 0.0635 - val_acc: 0.9798
Epoch 13/20
60000/60000 [=====] - 6s 96us/sample - loss:
0.0651 - acc: 0.9787 - val_loss: 0.0630 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 6s 96us/sample - loss:
0.0643 - acc: 0.9789 - val_loss: 0.0594 - val_acc: 0.9814
Epoch 15/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0600 - acc: 0.9807 - val_loss: 0.0601 - val_acc: 0.9818
Epoch 16/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0564 - acc: 0.9817 - val_loss: 0.0655 - val_acc: 0.9802
Epoch 17/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0573 - acc: 0.9810 - val_loss: 0.0623 - val_acc: 0.9808
Epoch 18/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0526 - acc: 0.9829 - val_loss: 0.0615 - val_acc: 0.9819
Epoch 19/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0471 - acc: 0.9841 - val_loss: 0.0613 - val_acc: 0.9811
Epoch 20/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0478 - acc: 0.9845 - val_loss: 0.0640 - val_acc: 0.9819
*****

```

Printing the Model Summary


Layer (type)	Output Shape	Param #
=====		

dense_30 (Dense)	(None, 512)	401920
batch_normalization_v1_6 (Ba	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_31 (Dense)	(None, 256)	131328
batch_normalization_v1_7 (Ba	(None, 256)	1024
dropout_7 (Dropout)	(None, 256)	0
dense_32 (Dense)	(None, 128)	32896
dense_33 (Dense)	(None, 10)	1290

Total params: 570,506
 Trainable params: 568,970
 Non-trainable params: 1,536

None

 10000/10000 [=====] - 1s 106us/sample - loss:
 0.0640 - acc: 0.9819
 Test score: 0.0640360280782188
 Test accuracy: 0.9819



```
In [33]: w_after = relumodel_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```



```
plt.xlabel('Output Layer ' )  
plt.show()
```



3 ReLU hidden Layers (512-128-64) + BatchNormalization + Dropout + ADAM

```
In [34]: relumodel_3 = tf.keras.models.Sequential()  
relumodel_3.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))  
relumodel_3.add(tf.keras.layers.BatchNormalization())  
relumodel_3.add(tf.keras.layers.Dropout(0.5))  
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))  
relumodel_3.add(tf.keras.layers.BatchNormalization())  
relumodel_3.add(tf.keras.layers.Dropout(0.5))  
relumodel_3.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))  
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))
```

```

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=b
atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 3,
                                     "Model": "3-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -
> 128 -> 64 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

```

```
vy = model.history['val_loss']  
ty = model.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 92us/sample - loss:
0.7178 - acc: 0.7768 - val_loss: 1.9599 - val_acc: 0.1297

Epoch 2/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.2933 - acc: 0.9117 - val_loss: 2.0435 - val_acc: 0.0982

Epoch 3/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.2226 - acc: 0.9337 - val_loss: 1.9565 - val_acc: 0.0989

Epoch 4/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.1820 - acc: 0.9456 - val_loss: 1.6727 - val_acc: 0.1564

Epoch 5/20

60000/60000 [=====] - 5s 78us/sample - loss:
0.1566 - acc: 0.9530 - val_loss: 1.3032 - val_acc: 0.3902

Epoch 6/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.1382 - acc: 0.9586 - val_loss: 0.7529 - val_acc: 0.7091

Epoch 7/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.1233 - acc: 0.9630 - val_loss: 0.3635 - val_acc: 0.9012

Epoch 8/20

60000/60000 [=====] - 5s 79us/sample - loss:
0.1130 - acc: 0.9647 - val_loss: 0.1758 - val_acc: 0.9527

Epoch 9/20

60000/60000 [=====] - 5s 78us/sample - loss:
0.1048 - acc: 0.9676 - val_loss: 0.1125 - val_acc: 0.9683

Epoch 10/20

60000/60000 [=====] - 5s 77us/sample - loss:
0.0934 - acc: 0.9713 - val_loss: 0.0882 - val_acc: 0.9721

Epoch 11/20

60000/60000 [=====] - 5s 78us/sample - loss:
0.0895 - acc: 0.9722 - val_loss: 0.0748 - val_acc: 0.9771

Epoch 12/20

60000/60000 [=====] - 5s 77us/sample - loss:

```

0.0841 - acc: 0.9740 - val_loss: 0.0722 - val_acc: 0.9784
Epoch 13/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0791 - acc: 0.9752 - val_loss: 0.0687 - val_acc: 0.9795
Epoch 14/20
60000/60000 [=====] - 5s 78us/sample - loss:
0.0737 - acc: 0.9768 - val_loss: 0.0689 - val_acc: 0.9784
Epoch 15/20
60000/60000 [=====] - 5s 79us/sample - loss:
0.0701 - acc: 0.9780 - val_loss: 0.0672 - val_acc: 0.9803
Epoch 16/20
60000/60000 [=====] - 5s 80us/sample - loss:
0.0672 - acc: 0.9791 - val_loss: 0.0656 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 5s 80us/sample - loss:
0.0621 - acc: 0.9801 - val_loss: 0.0690 - val_acc: 0.9805
Epoch 18/20
60000/60000 [=====] - 5s 79us/sample - loss:
0.0582 - acc: 0.9815 - val_loss: 0.0692 - val_acc: 0.9808
Epoch 19/20
60000/60000 [=====] - 5s 80us/sample - loss:
0.0567 - acc: 0.9818 - val_loss: 0.0654 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 5s 79us/sample - loss:
0.0520 - acc: 0.9833 - val_loss: 0.0659 - val_acc: 0.9812
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 512)	401920
batch_normalization_v1_8 (Ba	(None, 512)	2048
dropout_8 (Dropout)	(None, 512)	0
dense_35 (Dense)	(None, 128)	65664
batch_normalization_v1_9 (Ba	(None, 128)	512

dropout_9 (Dropout)	(None, 128)	0
dense_36 (Dense)	(None, 64)	8256
dense_37 (Dense)	(None, 10)	650
=====		
Total params: 479,050		
Trainable params: 477,770		
Non-trainable params: 1,280		


None

10000/10000 [=====] - 1s 102us/sample - loss:

0.0659 - acc: 0.9812

Test score: 0.06589906297894195

Test accuracy: 0.9812



```
In [35]: w_after = relumodel_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
```

```


out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



3 ReLU hidden Layers (384-256-128) + BatchNormalization + Dropout + ADAM

```
In [36]: relumodel_3 = tf.keras.models.Sequential()
relumodel_3.add(tf.keras.layers.Dense(384, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_3.add(tf.keras.layers.BatchNormalization())
relumodel_3.add(tf.keras.layers.Dropout(0.5))
relumodel_3.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.BatchNormalization())
relumodel_3.add(tf.keras.layers.Dropout(0.5))
relumodel_3.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_3.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_3.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_3.summary())
print("*****")

score = relumodel_3.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 3,
                                     "Model": "3-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 384 -> 256 -> 128 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": True,
```

```

        "Dropout-Present": True,
        "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
        "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
        "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
        "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])), ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 92us/sample - loss:
0.6575 - acc: 0.7929 - val_loss: 2.0845 - val_acc: 0.0976

Epoch 2/20

60000/60000 [=====] - 5s 77us/sample - loss:
0.2646 - acc: 0.9195 - val_loss: 2.5004 - val_acc: 0.0974

Epoch 3/20

60000/60000 [=====] - 5s 77us/sample - loss:
0.2092 - acc: 0.9355 - val_loss: 2.5218 - val_acc: 0.0975

Epoch 4/20

60000/60000 [=====] - 5s 77us/sample - loss:
0.1732 - acc: 0.9473 - val_loss: 2.5314 - val_acc: 0.0994

Epoch 5/20

60000/60000 [=====] - 5s 77us/sample - loss:
0.1480 - acc: 0.9542 - val_loss: 2.0204 - val_acc: 0.2185

Epoch 6/20

60000/60000 [=====] - 5s 76us/sample - loss:
0.1317 - acc: 0.9594 - val_loss: 1.4259 - val_acc: 0.3795


```
Epoch 7/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.1181 - acc: 0.9630 - val_loss: 0.6956 - val_acc: 0.7306
Epoch 8/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.1091 - acc: 0.9663 - val_loss: 0.2920 - val_acc: 0.9046
Epoch 9/20
60000/60000 [=====] - 5s 77us/sample - loss:
0.1014 - acc: 0.9682 - val_loss: 0.1586 - val_acc: 0.9481
Epoch 10/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0937 - acc: 0.9708 - val_loss: 0.0946 - val_acc: 0.9706
Epoch 11/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0886 - acc: 0.9713 - val_loss: 0.0796 - val_acc: 0.9744
Epoch 12/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0852 - acc: 0.9730 - val_loss: 0.0729 - val_acc: 0.9770
Epoch 13/20
60000/60000 [=====] - 5s 77us/sample - loss:
0.0769 - acc: 0.9758 - val_loss: 0.0673 - val_acc: 0.9784
Epoch 14/20
60000/60000 [=====] - 5s 77us/sample - loss:
0.0735 - acc: 0.9761 - val_loss: 0.0648 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 5s 92us/sample - loss:
0.0722 - acc: 0.9770 - val_loss: 0.0678 - val_acc: 0.9798
Epoch 16/20
60000/60000 [=====] - 5s 90us/sample - loss:
0.0663 - acc: 0.9779 - val_loss: 0.0633 - val_acc: 0.9799
Epoch 17/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0656 - acc: 0.9790 - val_loss: 0.0694 - val_acc: 0.9801
Epoch 18/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0599 - acc: 0.9808 - val_loss: 0.0641 - val_acc: 0.9805
Epoch 19/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0587 - acc: 0.9809 - val_loss: 0.0600 - val_acc: 0.9819
```


```
Epoch 20/20
60000/60000 [=====] - 5s 76us/sample - loss:
0.0557 - acc: 0.9824 - val_loss: 0.0643 - val_acc: 0.9819
*****
```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_38 (Dense)	(None, 384)	301440
batch_normalization_v1_10 (B	(None, 384)	1536
dropout_10 (Dropout)	(None, 384)	0
dense_39 (Dense)	(None, 256)	98560
batch_normalization_v1_11 (B	(None, 256)	1024
dropout_11 (Dropout)	(None, 256)	0
dense_40 (Dense)	(None, 128)	32896
dense_41 (Dense)	(None, 10)	1290
Total params: 436,746		
Trainable params: 435,466		
Non-trainable params: 1,280		

None

```
*****
10000/10000 [=====] - 1s 99us/sample - loss:
0.0643 - acc: 0.9819
Test score: 0.06429926064036845
Test accuracy: 0.9819
```



```
In [37]: w_after = relumodel_3.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ' )
plt.show()
```



5 Hidden Layers architecture

```
In [0]: final_output = final_output.append({"#Layers": "--",
                                             "Model": "--",
                                             "Layer-Architecture": "--",
                                             "Optimizer": "--", "BN-Present": "-
                                             ",
                                             "Dropout-Present": "--",
                                             "Train-loss": "--",
                                             "Test-loss": "--",
                                             "Train-accuracy": "--",
                                             "Test-Accuracy": "--"}, ignore_inde
                                             x=True)
```

5 ReLU hidden Layers (512-384-256-128-64) + ADAM

```
In [39]: relumodel_5 = tf.keras.models.Sequential()
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_5.add(tf.keras.layers.Dense(384, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 384 -> 256 -> 128 -> 64 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1])})
```

```

odel.history["acc"][n_epochs-1]), "Train-accuracy": '{:.5f}'.format(m
del.history["val_acc"][n_epochs-1])), "Test-Accuracy": '{:.5f}'.format(mo
ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 106us/sample - loss:
0.8376 - acc: 0.7474 - val_loss: 0.3146 - val_acc: 0.9047

Epoch 2/20

60000/60000 [=====] - 6s 93us/sample - loss:
0.2546 - acc: 0.9247 - val_loss: 0.2004 - val_acc: 0.9381

Epoch 3/20

60000/60000 [=====] - 6s 93us/sample - loss:
0.1716 - acc: 0.9494 - val_loss: 0.1515 - val_acc: 0.9550

Epoch 4/20

60000/60000 [=====] - 6s 94us/sample - loss:
0.1279 - acc: 0.9617 - val_loss: 0.1282 - val_acc: 0.9596

Epoch 5/20

60000/60000 [=====] - 6s 95us/sample - loss:
0.1027 - acc: 0.9691 - val_loss: 0.1003 - val_acc: 0.9683

Epoch 6/20

60000/60000 [=====] - 6s 94us/sample - loss:
0.0792 - acc: 0.9761 - val_loss: 0.0932 - val_acc: 0.9711

Epoch 7/20

60000/60000 [=====] - 6s 93us/sample - loss:
0.0665 - acc: 0.9795 - val_loss: 0.0875 - val_acc: 0.9739

Epoch 8/20

60000/60000 [=====] - 6s 94us/sample - loss:

```
0.0538 - acc: 0.9838 - val_loss: 0.0794 - val_acc: 0.9760
Epoch 9/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0431 - acc: 0.9869 - val_loss: 0.0839 - val_acc: 0.9742
Epoch 10/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0358 - acc: 0.9896 - val_loss: 0.0825 - val_acc: 0.9756
Epoch 11/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0305 - acc: 0.9909 - val_loss: 0.0763 - val_acc: 0.9772
Epoch 12/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0233 - acc: 0.9933 - val_loss: 0.0883 - val_acc: 0.9747
Epoch 13/20
60000/60000 [=====] - 6s 95us/sample - loss:
0.0219 - acc: 0.9936 - val_loss: 0.0822 - val_acc: 0.9763
Epoch 14/20
60000/60000 [=====] - 6s 93us/sample - loss:
0.0180 - acc: 0.9946 - val_loss: 0.0794 - val_acc: 0.9777
Epoch 15/20
60000/60000 [=====] - 6s 93us/sample - loss:
0.0134 - acc: 0.9963 - val_loss: 0.0816 - val_acc: 0.9766
Epoch 16/20
60000/60000 [=====] - 6s 93us/sample - loss:
0.0118 - acc: 0.9966 - val_loss: 0.0799 - val_acc: 0.9794
Epoch 17/20
60000/60000 [=====] - 6s 93us/sample - loss:
0.0093 - acc: 0.9976 - val_loss: 0.0833 - val_acc: 0.9795
Epoch 18/20
60000/60000 [=====] - 6s 93us/sample - loss:
0.0067 - acc: 0.9984 - val_loss: 0.0902 - val_acc: 0.9787
Epoch 19/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0048 - acc: 0.9989 - val_loss: 0.0876 - val_acc: 0.9793
Epoch 20/20
60000/60000 [=====] - 6s 94us/sample - loss:
0.0048 - acc: 0.9989 - val_loss: 0.0976 - val_acc: 0.9779
*****
Printing the Model Summary
```

Layer (type)	Output Shape	Param #
dense_42 (Dense)	(None, 512)	401920
dense_43 (Dense)	(None, 384)	196992
dense_44 (Dense)	(None, 256)	98560
dense_45 (Dense)	(None, 128)	32896
dense_46 (Dense)	(None, 64)	8256
dense_47 (Dense)	(None, 10)	650

Total params: 739,274
 Trainable params: 739,274
 Non-trainable params: 0


None

10000/10000 [=====] - 1s 123us/sample - loss:

0.0976 - acc: 0.9779

Test score: 0.09762186423194326

Test accuracy: 0.9779



```
In [40]: w_after = relumodel_5.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')
plt.show()
```



5 ReLU hidden Layers (512-256-128-64-32) + ADAM

```
In [41]: relumodel_5 = tf.keras.models.Sequential()
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_5.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 256 -> 128 -> 64 -> 32 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 87us/sample - loss:
0.9865 - acc: 0.7222 - val_loss: 0.3746 - val_acc: 0.8934
Epoch 2/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.3030 - acc: 0.9122 - val_loss: 0.2393 - val_acc: 0.9280
Epoch 3/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.2113 - acc: 0.9388 - val_loss: 0.1898 - val_acc: 0.9431
Epoch 4/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.1631 - acc: 0.9519 - val_loss: 0.1514 - val_acc: 0.9548
Epoch 5/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.1315 - acc: 0.9614 - val_loss: 0.1302 - val_acc: 0.9616
Epoch 6/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.1066 - acc: 0.9682 - val_loss: 0.1140 - val_acc: 0.9642
Epoch 7/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0913 - acc: 0.9722 - val_loss: 0.1091 - val_acc: 0.9658
Epoch 8/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0771 - acc: 0.9765 - val_loss: 0.0967 - val_acc: 0.9692
Epoch 9/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0653 - acc: 0.9811 - val_loss: 0.1044 - val_acc: 0.9686
Epoch 10/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0571 - acc: 0.9829 - val_loss: 0.0917 - val_acc: 0.9699
Epoch 11/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0499 - acc: 0.9851 - val_loss: 0.0842 - val_acc: 0.9751
Epoch 12/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0394 - acc: 0.9884 - val_loss: 0.0888 - val_acc: 0.9728
Epoch 13/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0358 - acc: 0.9896 - val_loss: 0.0878 - val_acc: 0.9726
```

```

Epoch 14/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0321 - acc: 0.9905 - val_loss: 0.0835 - val_acc: 0.9750
Epoch 15/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0250 - acc: 0.9931 - val_loss: 0.0862 - val_acc: 0.9752
Epoch 16/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0210 - acc: 0.9944 - val_loss: 0.0887 - val_acc: 0.9746
Epoch 17/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0192 - acc: 0.9949 - val_loss: 0.0850 - val_acc: 0.9761
Epoch 18/20
60000/60000 [=====] - 4s 73us/sample - loss:
0.0147 - acc: 0.9961 - val_loss: 0.0876 - val_acc: 0.9769
Epoch 19/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0122 - acc: 0.9970 - val_loss: 0.0819 - val_acc: 0.9778
Epoch 20/20
60000/60000 [=====] - 4s 74us/sample - loss:
0.0111 - acc: 0.9971 - val_loss: 0.0873 - val_acc: 0.9776
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_48 (Dense)	(None, 512)	401920
dense_49 (Dense)	(None, 256)	131328
dense_50 (Dense)	(None, 128)	32896
dense_51 (Dense)	(None, 64)	8256
dense_52 (Dense)	(None, 32)	2080
dense_53 (Dense)	(None, 10)	330
=====	=====	=====

Total params: 576,810

Trainable params: 576,810
Non-trainable params: 0


None

10000/10000 [=====] - 1s 97us/sample - loss:

0.0873 - acc: 0.9776

Test score: 0.08732793643142504

Test accuracy: 0.9776




```
In [42]: w_after = relu_model_5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5 ReLU hidden Layers (512-128-64-32-16) + ADAM

```
In [43]: relumodel_5 = tf.keras.models.Sequential()
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
```

```

relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(16, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.soft
max))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=b
atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -
> 128 -> 64 -> 32 -> 16 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
False,
                                     "Dropout-Present": False,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.
history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)

```



```

ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 74us/sample - loss: 1.4770 - acc: 0.5357 - val_loss: 0.7409 - val_acc: 0.7723

Epoch 2/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.4721 - acc: 0.8718 - val_loss: 0.3171 - val_acc: 0.9142

Epoch 3/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.2684 - acc: 0.9276 - val_loss: 0.2203 - val_acc: 0.9381

Epoch 4/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.1965 - acc: 0.9452 - val_loss: 0.1864 - val_acc: 0.9453

Epoch 5/20

60000/60000 [=====] - 4s 60us/sample - loss: 0.1565 - acc: 0.9559 - val_loss: 0.1476 - val_acc: 0.9586

Epoch 6/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.1301 - acc: 0.9627 - val_loss: 0.1297 - val_acc: 0.9622

Epoch 7/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.1091 - acc: 0.9683 - val_loss: 0.1208 - val_acc: 0.9665

Epoch 8/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.0940 - acc: 0.9732 - val_loss: 0.1123 - val_acc: 0.9678

Epoch 9/20

60000/60000 [=====] - 4s 61us/sample - loss: 0.0902 - acc: 0.9737 - val_loss: 0.1121 - val_acc: 0.9674

Epoch 10/20

60000/60000 [=====] - 4s 61us/sample - loss:

```

0.0728 - acc: 0.9786 - val_loss: 0.1048 - val_acc: 0.9708
Epoch 11/20
60000/60000 [=====] - 4s 62us/sample - loss:
0.0632 - acc: 0.9820 - val_loss: 0.1047 - val_acc: 0.9700
Epoch 12/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0547 - acc: 0.9844 - val_loss: 0.1009 - val_acc: 0.9716
Epoch 13/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0505 - acc: 0.9854 - val_loss: 0.0984 - val_acc: 0.9728
Epoch 14/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0422 - acc: 0.9878 - val_loss: 0.0943 - val_acc: 0.9722
Epoch 15/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0380 - acc: 0.9891 - val_loss: 0.0947 - val_acc: 0.9728
Epoch 16/20
60000/60000 [=====] - 4s 75us/sample - loss:
0.0338 - acc: 0.9905 - val_loss: 0.0961 - val_acc: 0.9731
Epoch 17/20
60000/60000 [=====] - 4s 66us/sample - loss:
0.0274 - acc: 0.9929 - val_loss: 0.1066 - val_acc: 0.9708
Epoch 18/20
60000/60000 [=====] - 4s 63us/sample - loss:
0.0249 - acc: 0.9934 - val_loss: 0.0993 - val_acc: 0.9724
Epoch 19/20
60000/60000 [=====] - 4s 61us/sample - loss:
0.0203 - acc: 0.9949 - val_loss: 0.1033 - val_acc: 0.9723
Epoch 20/20
60000/60000 [=====] - 4s 60us/sample - loss:
0.0186 - acc: 0.9952 - val_loss: 0.0968 - val_acc: 0.9739
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_54 (Dense)	(None, 512)	401920
dense_55 (Dense)	(None, 128)	65664

dense_56 (Dense)	(None, 64)	8256
dense_57 (Dense)	(None, 32)	2080
dense_58 (Dense)	(None, 16)	528
dense_59 (Dense)	(None, 10)	170
=====		
Total params: 478,618		
Trainable params: 478,618		
Non-trainable params: 0		


None

10000/10000 [=====] - 1s 84us/sample - loss:

0.0968 - acc: 0.9739

Test score: 0.09677388302332256

Test accuracy: 0.9739




```
In [44]: w_after = relumodel_5.get_weights()
         h1_w = w_after[0].flatten().reshape(-1,1)
```

```
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5 ReLU hidden Layers (512-384-256-128-64) + BatchNormalization + Dropout + ADAM

```
In [45]: relumodel_5 = tf.keras.models.Sequential()
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))
relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(384, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
```

```

relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.soft
max))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=b
atchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -
> 384 -> 256 -> 128 -> 64 -> 10",
                                     "Optimizer": "ADAM", "BN-Present":
True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model
.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model
.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(m
odel.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(mo
del.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

```

```
# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 154us/sample - loss: 1.1776 - acc: 0.6162 - val_loss: 2.1850 - val_acc: 0.1435

Epoch 2/20

60000/60000 [=====] - 8s 128us/sample - loss: 0.4080 - acc: 0.8788 - val_loss: 2.9801 - val_acc: 0.0974

Epoch 3/20

60000/60000 [=====] - 8s 129us/sample - loss: 0.3010 - acc: 0.9135 - val_loss: 2.9817 - val_acc: 0.0980

Epoch 4/20

60000/60000 [=====] - 8s 128us/sample - loss: 0.2378 - acc: 0.9312 - val_loss: 2.8179 - val_acc: 0.1386

Epoch 5/20

60000/60000 [=====] - 8s 127us/sample - loss: 0.2035 - acc: 0.9415 - val_loss: 1.9801 - val_acc: 0.3267

Epoch 6/20

60000/60000 [=====] - 8s 129us/sample - loss: 0.1769 - acc: 0.9487 - val_loss: 1.3237 - val_acc: 0.5080

Epoch 7/20

60000/60000 [=====] - 8s 128us/sample - loss: 0.1589 - acc: 0.9538 - val_loss: 0.7492 - val_acc: 0.7186

Epoch 8/20

60000/60000 [=====] - 8s 129us/sample - loss: 0.1448 - acc: 0.9576 - val_loss: 0.3424 - val_acc: 0.8903

Epoch 9/20

60000/60000 [=====] - 8s 130us/sample - loss: 0.1324 - acc: 0.9613 - val_loss: 0.1728 - val_acc: 0.9449

Epoch 10/20

60000/60000 [=====] - 8s 129us/sample - loss: 0.1248 - acc: 0.9650 - val_loss: 0.1008 - val_acc: 0.9679

Epoch 11/20

60000/60000 [=====] - 8s 130us/sample - loss:

```

0.1141 - acc: 0.9669 - val_loss: 0.0880 - val_acc: 0.9734
Epoch 12/20
60000/60000 [=====] - 8s 130us/sample - loss:
0.1042 - acc: 0.9695 - val_loss: 0.0800 - val_acc: 0.9776
Epoch 13/20
60000/60000 [=====] - 8s 130us/sample - loss:
0.1005 - acc: 0.9705 - val_loss: 0.0827 - val_acc: 0.9780
Epoch 14/20
60000/60000 [=====] - 8s 129us/sample - loss:
0.0926 - acc: 0.9731 - val_loss: 0.0796 - val_acc: 0.9782
Epoch 15/20
60000/60000 [=====] - 8s 129us/sample - loss:
0.0864 - acc: 0.9749 - val_loss: 0.0757 - val_acc: 0.9797
Epoch 16/20
60000/60000 [=====] - 8s 131us/sample - loss:
0.0821 - acc: 0.9757 - val_loss: 0.0744 - val_acc: 0.9804
Epoch 17/20
60000/60000 [=====] - 8s 130us/sample - loss:
0.0784 - acc: 0.9771 - val_loss: 0.0743 - val_acc: 0.9800
Epoch 18/20
60000/60000 [=====] - 8s 130us/sample - loss:
0.0768 - acc: 0.9775 - val_loss: 0.0695 - val_acc: 0.9823
Epoch 19/20
60000/60000 [=====] - 8s 129us/sample - loss:
0.0734 - acc: 0.9779 - val_loss: 0.0721 - val_acc: 0.9814
Epoch 20/20
60000/60000 [=====] - 8s 129us/sample - loss:
0.0680 - acc: 0.9805 - val_loss: 0.0734 - val_acc: 0.9817
*****

```

Printing the Model Summary


Layer (type)	Output Shape	Param #
dense_60 (Dense)	(None, 512)	401920
batch_normalization_v1_12 (B	(None, 512)	2048
dropout_12 (Dropout)	(None, 512)	0

dense_61 (Dense)	(None, 384)	196992
batch_normalization_v1_13 (B	(None, 384)	1536
dropout_13 (Dropout)	(None, 384)	0
dense_62 (Dense)	(None, 256)	98560
batch_normalization_v1_14 (B	(None, 256)	1024
dropout_14 (Dropout)	(None, 256)	0
dense_63 (Dense)	(None, 128)	32896
batch_normalization_v1_15 (B	(None, 128)	512
dropout_15 (Dropout)	(None, 128)	0
dense_64 (Dense)	(None, 64)	8256
dense_65 (Dense)	(None, 10)	650

```

Total params: 744,394
Trainable params: 741,834
Non-trainable params: 2,560
=====
None
*****
10000/10000 [=====] - 1s 148us/sample - loss:
0.0734 - acc: 0.9817
Test score: 0.07335510681418236
Test accuracy: 0.9817

```



```
In [46]: w_after = relumodel_5.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ' )  
plt.show()
```



5 ReLU hidden Layers (512-256-128-64-32) + BatchNormalization + Dropout + ADAM

```
In [47]: relumodel_5 = tf.keras.models.Sequential()  
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(256, activation=tf.nn.relu))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))
```

```

relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 256 -> 128 -> 64 -> 32 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)

```

```

ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 127us/sample - loss: 1.5776 - acc: 0.4854 - val_loss: 2.2071 - val_acc: 0.3203

Epoch 2/20

60000/60000 [=====] - 6s 102us/sample - loss: 0.6154 - acc: 0.8227 - val_loss: 2.5244 - val_acc: 0.1348

Epoch 3/20

60000/60000 [=====] - 6s 101us/sample - loss: 0.4099 - acc: 0.8871 - val_loss: 2.6993 - val_acc: 0.1823

Epoch 4/20

60000/60000 [=====] - 6s 101us/sample - loss: 0.3159 - acc: 0.9140 - val_loss: 2.4828 - val_acc: 0.2420

Epoch 5/20

60000/60000 [=====] - 6s 100us/sample - loss: 0.2600 - acc: 0.9298 - val_loss: 2.1246 - val_acc: 0.2918

Epoch 6/20

60000/60000 [=====] - 6s 100us/sample - loss: 0.2246 - acc: 0.9401 - val_loss: 1.2837 - val_acc: 0.5370

Epoch 7/20

60000/60000 [=====] - 6s 100us/sample - loss: 0.2021 - acc: 0.9465 - val_loss: 0.4334 - val_acc: 0.8613

Epoch 8/20

60000/60000 [=====] - 6s 101us/sample - loss: 0.1823 - acc: 0.9507 - val_loss: 0.2455 - val_acc: 0.9234

Epoch 9/20

60000/60000 [=====] - 6s 100us/sample - loss: 0.1640 - acc: 0.9564 - val_loss: 0.1540 - val_acc: 0.9529

Epoch 10/20

60000/60000 [=====] - 6s 101us/sample - loss:

```

0.1538 - acc: 0.9589 - val_loss: 0.1094 - val_acc: 0.9671
Epoch 11/20
60000/60000 [=====] - 6s 101us/sample - loss:
0.1414 - acc: 0.9630 - val_loss: 0.0929 - val_acc: 0.9744
Epoch 12/20
60000/60000 [=====] - 6s 100us/sample - loss:
0.1267 - acc: 0.9661 - val_loss: 0.0938 - val_acc: 0.9740
Epoch 13/20
60000/60000 [=====] - 6s 101us/sample - loss:
0.1230 - acc: 0.9676 - val_loss: 0.0906 - val_acc: 0.9761
Epoch 14/20
60000/60000 [=====] - 6s 102us/sample - loss:
0.1140 - acc: 0.9700 - val_loss: 0.0874 - val_acc: 0.9776
Epoch 15/20
60000/60000 [=====] - 6s 103us/sample - loss:
0.1078 - acc: 0.9718 - val_loss: 0.0839 - val_acc: 0.9793
Epoch 16/20
60000/60000 [=====] - 6s 102us/sample - loss:
0.1037 - acc: 0.9721 - val_loss: 0.0850 - val_acc: 0.9793
Epoch 17/20
60000/60000 [=====] - 6s 101us/sample - loss:
0.0976 - acc: 0.9740 - val_loss: 0.0839 - val_acc: 0.9780
Epoch 18/20
60000/60000 [=====] - 6s 102us/sample - loss:
0.0957 - acc: 0.9744 - val_loss: 0.0869 - val_acc: 0.9775
Epoch 19/20
60000/60000 [=====] - 6s 108us/sample - loss:
0.0908 - acc: 0.9761 - val_loss: 0.0902 - val_acc: 0.9788
Epoch 20/20
60000/60000 [=====] - 7s 110us/sample - loss:
0.0851 - acc: 0.9769 - val_loss: 0.0845 - val_acc: 0.9792
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_66 (Dense)	(None, 512)	401920
batch_normalization_v1_16 (B	(None, 512)	2048


dropout_16 (Dropout)	(None, 512)	0
dense_67 (Dense)	(None, 256)	131328
batch_normalization_v1_17 (Batch Normalization)	(None, 256)	1024
dropout_17 (Dropout)	(None, 256)	0
dense_68 (Dense)	(None, 128)	32896
batch_normalization_v1_18 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_69 (Dense)	(None, 64)	8256
batch_normalization_v1_19 (Batch Normalization)	(None, 64)	256
dropout_19 (Dropout)	(None, 64)	0
dense_70 (Dense)	(None, 32)	2080
dense_71 (Dense)	(None, 10)	330

=====

Total params: 580,650
Trainable params: 578,730
Non-trainable params: 1,920

None

10000/10000 [=====] - 2s 165us/sample - loss: 0.0845 - acc: 0.9792
Test score: 0.08452142000179738
Test accuracy: 0.9792



```
In [48]: w_after = relumodel_5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```



```
plt.xlabel('Output Layer ' )  
plt.show()
```



5 ReLU hidden Layers (512-128-64-32-16) + BatchNormalization + Dropout + ADAM

```
In [49]: relumodel_5 = tf.keras.models.Sequential()  
relumodel_5.add(tf.keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(input_dim, )))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(64, activation=tf.nn.relu))  
relumodel_5.add(tf.keras.layers.BatchNormalization())  
relumodel_5.add(tf.keras.layers.Dropout(0.5))  
relumodel_5.add(tf.keras.layers.Dense(32, activation=tf.nn.relu))
```

```

relumodel_5.add(tf.keras.layers.BatchNormalization())
relumodel_5.add(tf.keras.layers.Dropout(0.5))
relumodel_5.add(tf.keras.layers.Dense(16, activation=tf.nn.relu))
relumodel_5.add(tf.keras.layers.Dense(output_dim, activation=tf.nn.softmax))

relumodel_5.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model = relumodel_5.fit(x_train, y_train, epochs=n_epochs, batch_size=batchsize, verbose=1, validation_data=(x_test, y_test))

print("*****")
print("Printing the Model Summary")
print(relumodel_5.summary())
print("*****")

score = relumodel_5.evaluate(x_test, y_test)

print('Test score:', score[0])
print('Test accuracy:', score[1])

final_output = final_output.append({"#Layers": 5,
                                     "Model": "5-ReLU + Softmax",
                                     "Layer-Architecture": "784 -> 512 -> 128 -> 64 -> 32 -> 16 -> 10",
                                     "Optimizer": "ADAM", "BN-Present": True,
                                     "Dropout-Present": True,
                                     "Train-loss": '{:.5f}'.format(model.history["loss"][n_epochs-1]),
                                     "Test-loss": '{:.5f}'.format(model.history["val_loss"][n_epochs-1]),
                                     "Train-accuracy": '{:.5f}'.format(model.history["acc"][n_epochs-1]),
                                     "Test-Accuracy": '{:.5f}'.format(model.history["val_acc"][n_epochs-1])}, ignore_index=True)

fig,ax = plt.subplots(1,1)

```

```

ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,n_epochs+1))

vy = model.history['val_loss']
ty = model.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 111us/sample - loss: 1.8243 - acc: 0.3702 - val_loss: 2.1310 - val_acc: 0.2047

Epoch 2/20

60000/60000 [=====] - 5s 84us/sample - loss: 1.0923 - acc: 0.6401 - val_loss: 2.3769 - val_acc: 0.1413

Epoch 3/20

60000/60000 [=====] - 5s 84us/sample - loss: 0.7828 - acc: 0.7556 - val_loss: 2.3591 - val_acc: 0.1599

Epoch 4/20

60000/60000 [=====] - 5s 84us/sample - loss: 0.5977 - acc: 0.8222 - val_loss: 1.9045 - val_acc: 0.1981

Epoch 5/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.4813 - acc: 0.8649 - val_loss: 1.1301 - val_acc: 0.5712

Epoch 6/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.3978 - acc: 0.8934 - val_loss: 0.5505 - val_acc: 0.8466

Epoch 7/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.3391 - acc: 0.9116 - val_loss: 0.2975 - val_acc: 0.9232

Epoch 8/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.2975 - acc: 0.9240 - val_loss: 0.1832 - val_acc: 0.9515

Epoch 9/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.2686 - acc: 0.9325 - val_loss: 0.1413 - val_acc: 0.9633

Epoch 10/20

60000/60000 [=====] - 5s 83us/sample - loss:

```

0.2417 - acc: 0.9394 - val_loss: 0.1198 - val_acc: 0.9694
Epoch 11/20
60000/60000 [=====] - 5s 84us/sample - loss:
0.2292 - acc: 0.9432 - val_loss: 0.1151 - val_acc: 0.9714
Epoch 12/20
60000/60000 [=====] - 5s 85us/sample - loss:
0.2107 - acc: 0.9483 - val_loss: 0.1149 - val_acc: 0.9716
Epoch 13/20
60000/60000 [=====] - 5s 84us/sample - loss:
0.1958 - acc: 0.9520 - val_loss: 0.1137 - val_acc: 0.9733
Epoch 14/20
60000/60000 [=====] - 5s 83us/sample - loss:
0.1838 - acc: 0.9554 - val_loss: 0.1154 - val_acc: 0.9732
Epoch 15/20
60000/60000 [=====] - 5s 83us/sample - loss:
0.1755 - acc: 0.9575 - val_loss: 0.1127 - val_acc: 0.9753
Epoch 16/20
60000/60000 [=====] - 5s 83us/sample - loss:
0.1631 - acc: 0.9603 - val_loss: 0.1126 - val_acc: 0.9762
Epoch 17/20
60000/60000 [=====] - 5s 85us/sample - loss:
0.1568 - acc: 0.9626 - val_loss: 0.1140 - val_acc: 0.9764
Epoch 18/20
60000/60000 [=====] - 5s 85us/sample - loss:
0.1495 - acc: 0.9637 - val_loss: 0.1035 - val_acc: 0.9778
Epoch 19/20
60000/60000 [=====] - 5s 84us/sample - loss:
0.1400 - acc: 0.9659 - val_loss: 0.1041 - val_acc: 0.9778
Epoch 20/20
60000/60000 [=====] - 5s 85us/sample - loss:
0.1333 - acc: 0.9679 - val_loss: 0.0996 - val_acc: 0.9794
*****

```

Printing the Model Summary

Layer (type)	Output Shape	Param #
dense_72 (Dense)	(None, 512)	401920
batch_normalization_v1_20 (B	(None, 512)	2048


dropout_20 (Dropout)	(None, 512)	0
dense_73 (Dense)	(None, 128)	65664
batch_normalization_v1_21 (B	(None, 128)	512
dropout_21 (Dropout)	(None, 128)	0
dense_74 (Dense)	(None, 64)	8256
batch_normalization_v1_22 (B	(None, 64)	256
dropout_22 (Dropout)	(None, 64)	0
dense_75 (Dense)	(None, 32)	2080
batch_normalization_v1_23 (B	(None, 32)	128
dropout_23 (Dropout)	(None, 32)	0
dense_76 (Dense)	(None, 16)	528
dense_77 (Dense)	(None, 10)	170

=====

Total params: 481,562
Trainable params: 480,090
Non-trainable params: 1,472

None

10000/10000 [=====] - 1s 141us/sample - loss: 0.0996 - acc: 0.9794
Test score: 0.09960540832220577
Test accuracy: 0.9794



```
In [50]: w_after = relumodel_5.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10, 5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')\nplt.show()
```



Conclusion

Steps Summary -

- The dataset that is taken for analysing MLP techniques is MNIST which contains data about the handwritten images. Basically, we have to classify the handwritten numbers from the image to numeric.
- A number of architectures were deployed - **2 hidden layers, 3 hidden layers, 5 hidden layers**.
- In this, we specifically took Rectified Linear Unit (**ReLU**) as our default activation function with 'AdaM' optimizer.
- For each of the architecture, we tried BatchNormalization(**BN**) and **Dropout** (With dropout rate to 0.5) to see whether our model performs better or not.

```
In [51]: print("The below output summarizes the number of architecture performances that were deployed on MNIST Dataset")
print("*****")
final_output
```

The below output summarizes the number of architecture performances that were deployed on MNIST Dataset

Out[51]:

	#Layers	Model	Layer-Architecture	Optimizer	BN-Present	Dropout-Present	Train-loss	Test-loss	Train-accuracy	Accuracy
0	2	2-ReLU + Softmax	784 -> 512 -> 128 -> 10	ADAM	False	False	0.03062	0.06986	0.99197	(0.99197)
1	2	2-ReLU + Softmax	784 -> 256 -> 256 -> 10	ADAM	False	False	0.04022	0.07797	0.98873	(0.98873)
2	2	2-ReLU + Softmax	784 -> 384 -> 128 -> 10	ADAM	False	False	0.03553	0.07293	0.99040	(0.99040)
3	2	2-ReLU + Softmax	784 -> 512 -> 128 -> 10	ADAM	True	True	0.05164	0.05973	0.98323	(0.98323)
4	2	2-ReLU + Softmax	784 -> 256 -> 256 -> 10	ADAM	True	True	0.07511	0.06643	0.97567	(0.97567)
5	2	2-ReLU + Softmax	784 -> 384 -> 128 -> 10	ADAM	True	True	0.06306	0.06616	0.98013	(0.98013)
6	--	--	--	--	--	--	--	--	--	--
7	3	3-ReLU + Softmax	784 -> 512 -> 256 -> 128 -> 10	ADAM	False	False	0.01240	0.07288	0.99708	(0.99708)

	#Layers	Model	Layer-Architecture	Optimizer	BN-Present	Dropout-Present	Train-loss	Test-loss	Train-accuracy	Accuracy
8	3	3-ReLU + Softmax	784 -> 512 -> 128 -> 64 -> 10	ADAM	False	False	0.02145	0.07323	0.99483	(0.99483)
9	3	3-ReLU + Softmax	784 -> 384 -> 256 -> 128 -> 10	ADAM	False	False	0.01295	0.07172	0.99713	(0.99713)
10	3	3-ReLU + Softmax	784 -> 512 -> 256 -> 128 -> 10	ADAM	True	True	0.04785	0.06404	0.98450	(0.98450)
11	3	3-ReLU + Softmax	784 -> 512 -> 128 -> 64 -> 10	ADAM	True	True	0.05202	0.06590	0.98335	(0.98335)
12	3	3-ReLU + Softmax	784 -> 384 -> 256 -> 128 -> 10	ADAM	True	True	0.05572	0.06430	0.98243	(0.98243)
13	--	--	--	--	--	--	--	--	--	--
14	5	5-ReLU + Softmax	784 -> 512 -> 384 -> 256 -> 128 -> 64 -> 10	ADAM	False	False	0.00483	0.09762	0.99890	(0.99890)
15	5	5-ReLU + Softmax	784 -> 512 -> 256 -> 128 -> 64 -> 32 -> 10	ADAM	False	False	0.01107	0.08733	0.99713	(0.99713)
16	5	5-ReLU + Softmax	784 -> 512 -> 128 -> 64 -> 32 -> 16 -> 10	ADAM	False	False	0.01862	0.09677	0.99525	(0.99525)
17	5	5-ReLU + Softmax	784 -> 512 -> 384 -> 256 -> 128 -> 64 -> 10	ADAM	True	True	0.06797	0.07336	0.98052	(0.98052)
18	5	5-ReLU + Softmax	784 -> 512 -> 256 -> 128 -> 64 -> 32 -> 10	ADAM	True	True	0.08507	0.08452	0.97692	(0.97692)

#Layers	Model	Layer-Architecture	Optimizer	BN-Present	Dropout-Present	Train-loss	Test-loss	Train-accuracy	Ac
19	5	5-ReLU + Softmax 784 -> 512 -> 128 -> 64 -> 32 -> 16 -> 10	ADAM	True	True	0.13327	0.09961	0.96795	(

Output Representation -

- **Layers** -> The Number of Hidden Layers that were used.
- **Model** -> The Models that were used.
- **Layer Architecture** -> Number of neurons present in each hidden layer
- **Optimizer** -> The Type of Optimizer that was used.
- **BN-Present** -> Batch Normalization was used or not
- **Dropout-Present** -> Dropout rate was applied to the model or not

Output Conclusion -

By looking at the above table, We can conclude that after applying batch normalization and dropout, the loss(train & test) becomes significantly less different. Also, the accuracy for train as well test does not differs a lot.